

```

In [1]: import numpy as np
import matplotlib.pyplot as plt
import matplotlib.patches as patches
import seaborn as sns

class GridWorld:
    """
    GridWorld environment for TD Learning.
    Grid cell codes:
    0 = normal (step penalty)
    1 = goal (+1)
    2 = poison (-1)
    3 = wall (non-passable)
    """

    def __init__(self, grid_map, step_cost=-0.1, goal_reward=1.0, poison_penalty=-1):
        self.map = np.array(grid_map)
        self.num_rows, self.num_cols = self.map.shape
        self.num_states = self.num_rows * self.num_cols
        self.num_actions = 4 # up, right, down, left

        # Define rewards per cell type
        self.rewards = {
            0: step_cost,      # normal c
            1: goal_reward,    # goal
            2: poison_penalty, # poison
            3: 0.0             # wall
        }

        self.reward_function = self._build_reward_function()

    def reset(self):
        """
        Resets the environment to the starting state.
        Returns:
            state (int): Initial non-terminal, non-wall state
        """
        while True:
            r = np.random.randint(self.num_rows)
            c = np.random.randint(self.num_cols)
            if self.map[r, c] == 0: # Normal cell only
                self.current_state = self.get_state(r, c)
                break
        return self.current_state

    def step(self, action):
        """
        Takes an action from the current state and returns the result.
        Parameters:
            action (int): 0=up, 1=right, 2=down, 3=left
        Returns:
            next_state (int): resulting state
            reward (float): reward received
            done (bool): whether the episode has ended
        """

```

```

"""
r, c = self.get_position(self.current_state)

# Define motion directions
directions = [(-1, 0), (0, 1), (1, 0), (0, -1)]
dr, dc = directions[action]
new_r, new_c = r + dr, c + dc

# Check boundaries and walls
if 0 <= new_r < self.num_rows and 0 <= new_c < self.num_cols and self.map[new_r, new_c] != 2:
    next_state = self.get_state(new_r, new_c)
else:
    next_state = self.current_state # Bounce back

reward = self.reward_function[next_state]
done = self.map[new_r, new_c] in [1, 2] if (0 <= new_r < self.num_rows and
0 <= new_c < self.num_cols) else False

self.current_state = next_state
return next_state, reward, done

def _build_reward_function(self):
    rewards = np.zeros(self.num_states)
    for r in range(self.num_rows):
        for c in range(self.num_cols):
            s = self.get_state(r, c)
            cell_type = self.map[r, c]
            rewards[s] = self.rewards[cell_type]
    return rewards

def get_state(self, row, col):
    return row * self.num_cols + col

def get_position(self, state):
    return divmod(state, self.num_cols)

def display_map(self):
    """
    Displays the GridWorld layout (walls, goal, poison, etc.)
    """
    fig, ax = plt.subplots(figsize=(self.num_cols, self.num_rows))
    cmap = {0: 'white', 1: '#00917C', 2: '#FF5252', 3: 'black'}

    for r in range(self.num_rows):
        for c in range(self.num_cols):
            cell_type = self.map[r, c]
            rect = patches.Rectangle((c, self.num_rows - r - 1), 1, 1,
                                     facecolor=cmap[cell_type], edgecolor='gray')
            ax.add_patch(rect)

    ax.set_xlim(0, self.num_cols)
    ax.set_ylim(0, self.num_rows)
    ax.set_xticks([])
    ax.set_yticks([])
    ax.set_title("GridWorld Layout")
    plt.show()

```

```

def show_state_numbering(self):
    """
    Displays the grid with state numbers annotated in each cell.
    """
    fig, ax = plt.subplots(figsize=(self.num_cols, self.num_rows))
    cmap = {0: 'white', 1: '#00917C', 2: '#FF5252', 3: 'black'}

    for r in range(self.num_rows):
        for c in range(self.num_cols):
            s = self.get_state(r, c)
            cell_type = self.map[r, c]
            color = cmap[cell_type]

            rect = patches.Rectangle((c, self.num_rows - r - 1), 1, 1,
                                     facecolor=color, edgecolor='gray')
            ax.add_patch(rect)

            if cell_type != 3:
                ax.text(c + 0.5, self.num_rows - r - 0.5,
                        f"s = {s}", ha='center', va='center', fontsize=10, weight='bold')

    ax.set_xlim(0, self.num_cols)
    ax.set_ylim(0, self.num_rows)
    ax.set_aspect('equal')
    ax.axis('off')
    ax.set_title("GridWorld State Numbering", fontsize=14)
    plt.show()

def display_reward_map(self):
    """
    Displays the grid with reward values annotated for each state.
    """
    fig, ax = plt.subplots(figsize=(self.num_cols, self.num_rows))
    cmap = {0: 'white', 1: '#00917C', 2: '#FF5252', 3: 'black'}

    for r in range(self.num_rows):
        for c in range(self.num_cols):
            cell_type = self.map[r, c]
            s = self.get_state(r, c)
            color = cmap[cell_type]

            rect = patches.Rectangle((c, self.num_rows - r - 1), 1, 1,
                                     facecolor=color, edgecolor='gray')
            ax.add_patch(rect)

            if cell_type != 3:
                reward = self.reward_function[s]
                ax.text(c + 0.5, self.num_rows - r - 0.5,
                        f"R = {reward:.2f}", ha='center', va='center', fontsize=10, weight='bold')

    ax.set_xlim(0, self.num_cols)
    ax.set_ylim(0, self.num_rows)
    ax.set_aspect('equal')
    ax.axis('off')
    ax.set_title("Reward Function Map")
    plt.show()

```

```

def plot_state_values(self, values):
    """
    Plot state values as a heatmap.
    """
    reshaped = values.reshape(self.num_rows, self.num_cols)
    plt.figure(figsize=(self.num_cols+1.25, self.num_rows+1.25))
    ax = sns.heatmap(reshaped, annot=True, fmt=".2f", cmap="coolwarm",
                     annot_kws={"size": 14}, square=True, linewidths=0.5)
    ax.set_xticks([])
    ax.set_yticks([])
    plt.show()

def _quatromatrix(self, left, bottom, right, top, ax=None, triplotkw={}, tripco
if not ax:
    ax = plt.gca()
n, m = left.shape

a = np.array([[0, 0], [0, 1], [0.5, 0.5], [1, 0], [1, 1]])
tr = np.array([[0, 1, 2], [0, 2, 3], [2, 3, 4], [1, 2, 4]])

A = np.zeros((n * m * 5, 2))
Tr = np.zeros((n * m * 4, 3), dtype=int)

for i in range(n):
    for j in range(m):
        k = i * m + j
        A[k * 5:(k + 1) * 5, :] = np.c_[a[:, 0] + j, a[:, 1] + i]
        Tr[k * 4:(k + 1) * 4, :] = tr + k * 5

C = np.c_[left.flatten(), bottom.flatten(), right.flatten(), top.flatten()]

ax.triplot(A[:, 0], A[:, 1], Tr, **triplotkw)
tripcolor = ax.tripcolor(A[:, 0], A[:, 1], Tr, facecolors=C, **tripcolorkw)
return tripcolor

def plot_action_values(self, q_values):
    """
    Visualizes Q-values (action-values) for each state in the grid.
    Triangles in each cell indicate value of Up (0), Right (1), Down (2), Left
    """
    num_states, num_actions = q_values.shape
    assert num_states == self.num_states and num_actions == self.num_actions

    rows, cols = self.num_rows, self.num_cols
    top = q_values[:, 0].reshape((rows, cols))
    right = q_values[:, 1].reshape((rows, cols))
    bottom = q_values[:, 2].reshape((rows, cols))
    left = q_values[:, 3].reshape((rows, cols))

    # Text annotation positions
    top_pos = [(j + 0.38, i + 0.25) for i in range(rows) for j in range(cols)]
    right_pos = [(j + 0.65, i + 0.5) for i in range(rows) for j in range(cols)]
    bottom_pos = [(j + 0.38, i + 0.8) for i in range(rows) for j in range(cols)]
    left_pos = [(j + 0.05, i + 0.5) for i in range(rows) for j in range(cols)]

```

```

fig, ax = plt.subplots(figsize=(cols*2.0, rows*2.0))
ax.set_ylim(rows, 0)

# Triangular heatmap with Q-values
tripcolor = self._quatromatrix(left, bottom, right, top, ax=ax,
                               triplotkw={"color": "k", "lw": 1},
                               tripcolorkw={"cmap": "coolwarm"})

# Add text annotations
for i, (x, y) in enumerate(top_pos):
    ax.text(x, y, f"{top.flatten()[i]:.2f}", size=11, color="w")
for i, (x, y) in enumerate(right_pos):
    ax.text(x, y, f"{right.flatten()[i]:.2f}", size=11, color="w")
for i, (x, y) in enumerate(bottom_pos):
    ax.text(x, y, f"{bottom.flatten()[i]:.2f}", size=11, color="w")
for i, (x, y) in enumerate(left_pos):
    ax.text(x, y, f"{left.flatten()[i]:.2f}", size=11, color="w")

ax.margins(0)
ax.set_aspect("equal")
fig.colorbar(tripcolor)
ax.set_title("Action-Value Function (Q-values)")
plt.show()

```

```

In [2]: # Define a 3x4 grid map Layout:
# 0 = normal cell, 1 = goal, 2 = poison, 3 = wall

grid_map = [
    [0, 0, 0, 1],
    [0, 3, 0, 2],
    [0, 0, 0, 0]
]

# Initialize the GridWorld environment
env = GridWorld(grid_map)

```

```

In [3]: # Show the grid with color-coded cells
env.display_map()

```

GridWorld Layout


```
In [4]: # Call the function to display state numbering
env.show_state_numbering()
```

GridWorld State Numbering

<b>s = 0</b>	<b>s = 1</b>	<b>s = 2</b>	<b>s = 3</b>
<b>s = 4</b>		<b>s = 6</b>	<b>s = 7</b>
<b>s = 8</b>	<b>s = 9</b>	<b>s = 10</b>	<b>s = 11</b>

```
In [5]: # Call the function to display rewards
env.display_reward_map()
```

Reward Function Map

R = -0.10	R = -0.10	R = -0.10	R = 1.00
R = -0.10		R = -0.10	R = -1.00
R = -0.10	R = -0.10	R = -0.10	R = -0.10

```

In [6]: def td_zero_prediction(env, num_episodes=5000, discount_factor=0.99, learning_rate=
        """
        Implements the TD(0) prediction algorithm to estimate state values.

        Parameters:
        - env: GridWorld environment object
        - num_episodes: number of episodes to run
        - discount_factor (gamma): how much future rewards are valued
        - learning_rate (alpha): step size for updating value estimates

        Returns:
        - V: A NumPy array of shape [num_states], with estimated state values
        """
        V = np.zeros(env.num_states) # Initialize state-value function

        # Set terminal states (goal and poison) to their known final rewards
        for s in range(env.num_states):
            r, c = env.get_position(s)
            if env.map[r, c] in [1, 2]: # Goal or poison
                V[s] = env.reward_function[s]

        for episode in range(num_episodes):
            state = env.reset()

            while True:
                # Choose a random action (uniform policy)
                action = np.random.randint(env.num_actions)

                # Take a step in the environment
                next_state, reward, done = env.step(action)

                # TD(0) Update
                V[state] += learning_rate * (
                    reward + discount_factor * V[next_state] - V[state]
                )

                # Move to next state
                state = next_state

                if done:
                    break

        return V


```

```

In [7]: # Run TD(0)
td_values = td_zero_prediction(env, num_episodes=5000, discount_factor=0.99, learni

# Visualize the learned value function
env.plot_state_values(td_values)

```



```
In [8]: import numpy as np
import random

def sarsa(env, num_episodes=5000, alpha=0.01, gamma=0.99, epsilon_decay=0.99, min_e
"""
SARSA algorithm: On-policy Temporal-Difference control to estimate the Q-functi
"""

# Initialize Q-table with zeros: shape = [number of states x number of actions]
Q = np.zeros((env.num_states, env.num_actions))

# Optionally initialize Q-values for terminal states for stability
for s in range(env.num_states):
    r, c = env.get_position(s) # Get (row, col) of the state
    if env.map[r, c] == 1:     # Goal state
        Q[s, :] = 1.0
    elif env.map[r, c] == 2:   # Trap (death) state
        Q[s, :] = -1.0

# Define epsilon-greedy action selection strategy
def get_action(q_values, epsilon):
    if random.random() < epsilon: # With probability ε, explore
        return random.randint(0, env.num_actions - 1)
    else:
        # Otherwise, exploit the best action
        return np.argmax(q_values)

epsilon = 1.0 # Start with full exploration

# Loop over episodes
for episode in range(num_episodes):
    state = env.reset() # Reset environment to a random starting state
    action = get_action(Q[state], epsilon) # Choose initial action
    done = False # Track episode termination
```



```

while not done:
    next_state, reward, done = env.step(action) # Take action → observe ou
    next_action = get_action(Q[next_state], epsilon) # Choose next action

    # Compute the TD target and error
    td_target = reward + gamma * Q[next_state][next_action]
    td_error = td_target - Q[state][action]

    # Update the Q-value towards the TD target
    Q[state][action] += alpha * td_error

    # Move to the next state and action
    state = next_state
    action = next_action

    # Decay ε after each episode, but keep it above a minimum value
    epsilon = max(min_epsilon, epsilon * epsilon_decay)

return Q # Return the Learned Q-table


```

```

In [9]: # Run SARSA
q_sarsa = sarsa(env, num_episodes=5000)

# Visualize the Q-values
env.plot_action_values(q_sarsa)

```



```

In [10]: import numpy as np
import random

```

```

def q_learning(env, num_episodes=5000, alpha=0.01, gamma=0.99, epsilon_decay=0.99,
    """
    Q-Learning algorithm: Off-policy TD control to estimate the optimal Q-function.
    """

    # Initialize Q-table with zeros: shape = [number of states x number of actions]
    Q = np.zeros((env.num_states, env.num_actions))

    # Optionally initialize Q-values for terminal states for faster convergence
    for s in range(env.num_states):
        r, c = env.get_position(s) # Get (row, col) position
        if env.map[r, c] == 1:      # Goal state
            Q[s, :] = 1.0
        elif env.map[r, c] == 2:    # Trap state
            Q[s, :] = -1.0

    # Define  $\epsilon$ -greedy policy for action selection
    def get_action(q_values, epsilon):
        if random.random() < epsilon: # Explore
            return random.randint(0, env.num_actions - 1)
        else:                          # Exploit
            return np.argmax(q_values)

    epsilon = 1.0 # Initial exploration probability

    # Loop through all episodes
    for episode in range(num_episodes):
        state = env.reset() # Start from a random non-terminal state
        done = False # Whether the episode has ended

        while not done:
            # Choose action using current  $\epsilon$ -greedy policy
            action = get_action(Q[state], epsilon)

            # Perform the action and observe outcome
            next_state, reward, done = env.step(action)

            # Q-Learning update uses the maximum Q-value of next state (greedy)
            best_next_q = np.max(Q[next_state])
            td_target = reward + gamma * best_next_q # TD target
            td_error = td_target - Q[state][action] # TD error

            # Update the Q-value
            Q[state][action] += alpha * td_error

            # Transition to the next state
            state = next_state

        # Decay  $\epsilon$  at the end of each episode
        epsilon = max(min_epsilon, epsilon * epsilon_decay)

    return Q # Return the Learned Q-table


```

```

In [11]: # Run Q-Learning
q_learning = q_learning(env, num_episodes=5000)

```

```
# Visualize the Q-values
env.plot_action_values(q_learning)
```



In [ ]: